

Deep Learning for Speech and Language Processing

Exercise Sheet 3: Backpropagation, Gradient Descent

Maximilian Schmidt, Dirk Väth, Ngoc Thang Vu

18th of November 2019

Deadline & Submission

Please submit your solutions until Monday, November 25th, 2019, 5:30pm. You may also see the deadline at any time on the website under *Exercises → Upload*.

Submit the results for the calculation task to the website. Your solutions for the theory task should be submitted to ILIAS.

Notation

Additional notation conventions for this exercise:

Symbol	Example	Description
d	$\frac{df(z)}{dz}$	full derivative of f (f only depends on z)
∂	$\frac{\partial f}{\partial z}$	partial derivative of f with respect to z (f might depend on other inputs)
Greek delta	$\delta^l = \frac{\partial C}{\partial z^l}$	partial derivative of the cost function C with respect to the inputs to the activation function in layer l of a neural network

Theory Tasks

Theory Task 1: Activation derivatives

In the lecture you were shown the tanh function as alternative activation function in neural networks.

Exercise 1.

Compute the first derivative of $\tanh(z) = \frac{\sinh(z)}{\cosh(z)} = \frac{e^z - e^{-z}}{e^z + e^{-z}}$ with respect to z .

Hint: $\frac{d\sinh(z)}{dz} = \cosh(z)$, $\frac{d\cosh(z)}{dz} = \sinh(z)$. You can either first apply the product rule and then the chain rule or directly work with the quotient rule $f(x) = \frac{g(x)}{h(x)}$; $f'(x) = \frac{g'(x)h(x) - g(x)h'(x)}{[h(x)]^2}$.

Theory Task 2: Mean Squared Error vs Cross-Entropy Cost Function

In this exercise, we will look at the effect of the choice of the cost function on the gradients. Consider the network for multi-class classification with $N = 5$ classes, which only consists of an output layer, with f being the softmax function, depicted in Figure 1.

In order to train the parameters of this network, i.e. W and b , we need a cost function, which evaluates the predictions of the network y against a ground truth \hat{y} . You already know two cost functions for neural networks:

Mean Squared Error (MSE) $c_{\text{MSE}}(\hat{y}, y) = \sum_{i=0}^{N-1} (\hat{y}_i - y_i)^2$

Cross Entropy (CE) $c_{\text{CE}}(\hat{y}, y) = -\sum_{i=1}^N \hat{y}_i \ln y_i = -\ln y_t$ for $\hat{y}_t = 1$ in the case of classification




Figure 1: Network for multi-class classification.

Exercise 2.

- (1) Give a general formula on how to compute the error c of the network given an input x . Then using the chain rule, derive the general formula to compute the derivative of c with respect to a weight w_{ij} , which connects input x_j to z_i . Hint: Use the derivative of the softmax function: $\frac{\partial y_i}{\partial z_j} = y_j(1 - y_j)$ for $i = j$ and $= -y_i y_j$ for $i \neq j$.
- (2) What do you need to change in a) to compute the derivative of c with respect to a bias entry b_i ?
- (3) Compute the derivative of c_{CE} with respect to a weight w_{ij} .
- (4) Compute the derivative of $\frac{\partial c_{MSE}}{\partial w_{ij}}$. What happens with the derivative for values of y_i close to 0 or 1?
- (5) Why is it important to shuffle the dataset after each training pass through the dataset (= epoch) when training with stochastic or mini-batch gradient descent?

Calculation Tasks

Stochastic Gradient Descent

In this exercise, we will compute a complete step of the stochastic gradient descent algorithm by hand. This includes a forward pass to compute the output and backward pass to compute the gradients. We will use the multi-class banana network from exercise sheet 2 with some modifications as depicted in Figure 2. The task of the network is

still to predict $y = \begin{bmatrix} p(\text{banana}) \\ p(\text{cookie}) \\ p(\text{other}) \end{bmatrix}$ given a tweet in bag-of-words representation as input.

The network has non-linear activation in the two hidden layers, specifically $f^1 = f^2 = \text{sigmoid} = \sigma$. The output layer uses softmax activation. There is only a bias vector in the first hidden layer initialized to 0. The second hidden layer and the output layer do not have biases. The network has the following initial weight matrices:

$$W^1 \in \mathbb{R}^{2 \times 4} = \begin{bmatrix} 1 & 0 & 1 & 1 \\ 0.5 & 0 & 0.5 & 0.5 \end{bmatrix}, W^2 \in \mathbb{R}^{2 \times 2} = \begin{bmatrix} 0 & -1 \\ 0 & 1 \end{bmatrix}, W^3 \in \mathbb{R}^{3 \times 2} = \begin{bmatrix} 0.5 & 1 \\ 0.5 & 0.5 \\ 0 & 0 \end{bmatrix}.$$

We will refer to the set of all parameters of the network with $\theta = \{W^1, W^2, W^3, b^1\}$. In the forward pass, we will first compute the output of the network given some input and score the network output using CE. Subsequently, in the backward pass, we will compute the gradients of the parameters $\Delta c_{CE}(\theta)$ with respect to the cost function and make a gradient update such that the likelihood of the ground truth increases: $\theta^1 \leftarrow \theta - \eta \nabla c_{CE}(\theta)$.

Exercise 3.

- (1) Forward pass: Compute the output y of the network for the tweet ‘A banana a day keeps the doctor away.’ Represent the tweet as a 4-dimensional bag-of-words vectors given the 4-word vocabulary

$$V = \{v_0 : \text{sweet}, v_1 : \text{banana}, v_2 : \text{cookie}, v_3 : \text{chocolate}\}.$$

Hint: Write down the intermediate results a^l, z^l as you will need them in Question (3).

- (2) Compute the cost c_{CE} (using ln) for the output y , given the correct label $\hat{y} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$.

- (3) Backward pass: Compute the first derivatives for the parameters W_1 of the banana network. For consistency with the lecture, use the denominator layout to represent the gradients, i.e. $\frac{\partial c_{CE}}{W}$ for $W \in \mathbb{R}^{n \times m}$ also has $n \times m$ dimensions.

Hint: You need to successively compute the vectors δ^l in each layer, starting with the output layer. Then you can obtain $\frac{\partial c_{CE}}{w_{ij}^l} = \frac{\partial c_{CE}}{\partial z_i^l} \frac{\partial z_i^l}{\partial w_{ij}^l} = \delta_i^l a_j^{l-1}$. The computation for the $\frac{\partial c_{CE}}{b_j^1}$ works analogously with a small change. You will need all derivatives in the following tasks.

- (4) For the parameters W^1

(1) compute the first derivatives.

(2) perform a gradient update step using the derivatives you computed in task 3.4.1 with $\eta = 0.1$ (provide the updated weight matrix).

- (5) For the parameters W^2

(1) compute the first derivatives.

(2) perform a gradient update step using the derivatives you computed in task 3.5.1 with $\eta = 0.1$ (provide the updated weight matrix).

- (6) For the parameters W^3

(1) compute the first derivatives.

(2) perform a gradient update step using the derivatives you computed in task 3.6.1 with $\eta = 0.1$ (provide the updated weight matrix).

- (7) For the parameters b^1

(1) compute the first derivatives.

(2) perform a gradient update step using the derivatives you computed in task 3.7.1 with $\eta = 0.1$ (provide the updated weight matrix).




Figure 2: Revised banana network.