

# Deep Learning for Speech and Language Processing

## Exercise Sheet 6: RNNs, LSTMs

Maximilian Schmidt, Pascal Tilli, Dirk Väth, Ngoc Thang Vu

December 14th, 2020

### Deadline

Please submit your solutions until Monday, December 21st, 2020, noon. You may also see the deadline at any time on the website under *Exercises → Upload*.

### Notation

Notation conventions for this exercise:

Symbol	Description
$[x, y]$	concatenation of two vectors: $z = [x, y]$ for $x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \in \mathbb{R}^2, y = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix} \in \mathbb{R}^3 \rightarrow z = \begin{bmatrix} x_1 \\ x_2 \\ y_1 \\ y_2 \\ y_3 \end{bmatrix} \in \mathbb{R}^{2+3}$
$x_t$	input <i>vector</i> at timestep/position $t$ in a sequence of vectors (e.g. RNN input/output)
$x_{t_i}$	$i$ -th entry of a vector (= scalar) at timestep $t$ of a sequence
$\odot$	element-wise product of two vectors

### Submission

#### Calculation Tasks

Upload (only solutions!) to <https://dlcourse.ims.uni-stuttgart.de/> using your account.

#### Theory Tasks

You don't have to submit this part, but it will help you with the contents of this course.

### Warm-Up

#### Exercise 1.

- (1) **Calculation Task** Compute the number of parameters with input layer size 100, hidden layer size 400 and output layer size 5 for
  - (1) an Elman RNN.
  - (2) a Jordan RNN.

Always include the bias!

- (2) **Theory Task** In an Elman network, the output of the hidden layer at timestep  $t$   $a_t$  is stored in the memory and applied to compute the output of the hidden layer in the next timestep  $t + 1$ . However, in a Jordan network, the *network's output*  $y_t$  is stored in the memory and used to compute the output of the hidden layer

at  $t + 1$ . Give the equations to compute  $a_t$  for both network types. What is the consequence on the hidden and output layer sizes of Elman vs. Jordan networks?

(3) **Theory Task** The computations in an LSTM cell can be defined by the following equations:

$$f_t = \sigma(W_f[h_{t-1}, x_t] + b_f) \quad (\sigma = \text{logistic sigmoid}) \quad (1)$$

$$i_t = \sigma(W_i[h_{t-1}, x_t] + b_i) \quad (2)$$

$$\tilde{C}_t = \tanh(W_C[h_{t-1}, x_t] + b_C) \quad (3)$$

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t \quad (4)$$

$$o_t = \sigma(W_o[h_{t-1}, x_t] + b_o) \quad (5)$$

$$h_t = o_t \odot \tanh(C_t) \quad (6)$$

In what range are the entries of the input, forget and output gate? In what range are the entries of the candidate activation  $\tilde{C}_t$ ? And in what range are the outputs of the LSTM cell?

## Theory Task: RNN

In this exercise we compute the output of an Elman RNN for POS tagging, which is depicted in Figure 1. Use the logistic sigmoid as the activation function  $\sigma$ .

### Exercise 2.

Compute the sequence of POS tags for the input sentence ‘Kevin likes bananas’. Represent each word as one-hot vector given the 3-word vocabulary  $V = \{v_1 : \text{Kevin}, v_2 : \text{likes}, v_3 : \text{bananas}\}$ .

The outputs  $y_t$  of the network in each timestep are 3-dimensional vectors  $\begin{bmatrix} p(\text{VBZ}) \\ p(\text{NNS}) \\ p(\text{NNP}) \end{bmatrix}$  denoting the probability of the following POS tags according to the Penn Treebank Tagset: verb in 3rd person singular present, plural noun, proper noun singular. The memory ( $a_0$ ) is initialized as zero vector. Use the following weights of the network, there are no biases in any layer.

$$W_i = \begin{bmatrix} 1 & -1 & 0 \\ -1 & 0 & 0.5 \end{bmatrix}, W_h = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}, W_{out} = \begin{bmatrix} -1 & 1 \\ 0 & -1 \\ 1 & 0 \end{bmatrix}.$$

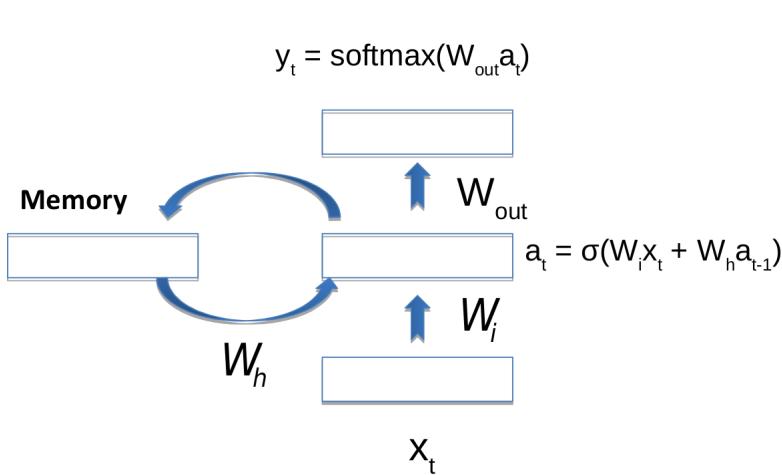


Figure 1: RNN architecture for Exercise 2.

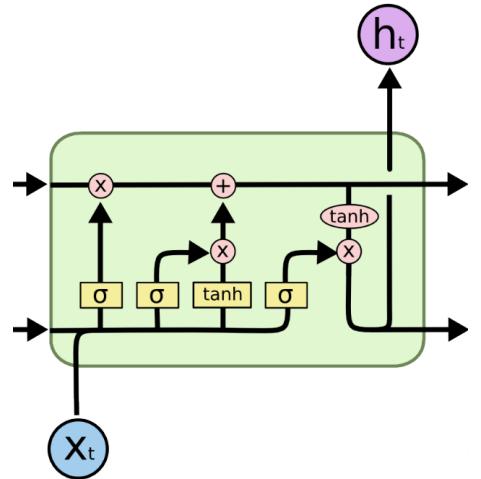


Figure 2: LSTM cell  
(<http://colah.github.io/posts/2015-08-Understanding-LSTMs>).

## Calculation Task: LSTM

To have more control over the memory updates in RNNs, we can replace the simple RNN cell, in the RNN architecture of Exercise 2 with an LSTM cell as depicted in Figure 2. By doing so the memory vectors  $a_t$  become

tuples of a hidden state and cell state  $(h_t, C_t)$ , which are computed according to Equation 1-6 given in Exercise 1.3. Note that in LSTMs only the hidden state of the memory tuple is passed on to the output layer.

**Exercise 3.**

Determine the predicted POS tag for the second word of the input sentence from Exercise 2. Assume as hidden state after processing the first word  $h_1 = \begin{bmatrix} 0.370 \\ -0.123 \end{bmatrix}, C_1 = \begin{bmatrix} 0.557 \\ -0.338 \end{bmatrix}$ . The weight matrices of the LSTM cell and for the output are given below, all bias vectors are zero vectors.

$$W_f = \begin{bmatrix} 0 & -1 & -1 & 1 & 0 \\ -1 & 1 & 0 & -1 & 1 \end{bmatrix}, W_i = \begin{bmatrix} -1 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & -1 \end{bmatrix}, W_C = \begin{bmatrix} -1 & 0 & 1 & 1 & 0 \\ -1 & -1 & 1 & -1 & 1 \end{bmatrix}$$

$$W_o = \begin{bmatrix} -1 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & -1 & 1 \end{bmatrix}, W_{out} = \begin{bmatrix} -1 & 1 \\ 0 & -1 \\ 0 & 0 \end{bmatrix}$$